

Questions & Confirmations

A short list of things I'd like you to confirm so I can set the system up correctly for the full source list. Most can be answered in a sentence or a yes/no.

1 Are the source links final?

In last quarter's set, some links opened a web page that had a download button, and the button was what actually took you to the PDF. Going forward, not just for these 37 but for future batches too, can I assume the links will be curated by hand so each one is final? So if a link is a web page, I'll treat that page as the thing to read, and if it's a PDF, the link goes straight to the file. Or is that extra click something I should plan to automate down the line?

2 Are there sources that should be read together?

Last quarter, a few calls came from two documents combined, where the view only really made sense once you read both together (say a review document alongside an outlook from the same firm). Will that come up again in the 37-source list? If so, it'd help to have a quick list of which titles go together: just the firm and the two titles, like *"J.P. Morgan: 'Title A' + 'Title B', read together."* A short note or a few lines in an email is plenty.

3 Can you confirm the call legend?

Quick one to confirm: I'm recording each call as O for overweight, N for neutral, U for underweight, and UNCERTAIN when the document doesn't really commit either way. Does that match what you'd like to see in the output?

4 Does a small change in a forecast number count as a call?

Sometimes a document only nudges a forecast number rather than stating an actual view. Should that count as a call? For now I've assumed a move needs to be at least around 25 basis points for rates, or 2% for currencies, before it's worth recording. Is that a reasonable line to draw, or would you put it somewhere else?

5 Should the system read between the lines, or only record stated views?

Some views are implied rather than spelled out, like a broad economic comment that clearly points to a call without saying it outright. Do you want the system to make those judgement calls, or just stick to what's explicitly stated? Right now it does include the implied ones, but marks them lower-confidence and flags them for review.

6 How specific should the calls be?

If a document names a specific country (say, Taiwan equities), should I record the call at the country level, or roll it up to the region (Asia)? More generally, when the wording sits between a broad category and a specific one, which do you prefer I record?

7 For dial or scale-style pages, which levels do I record?

Some documents present their views as dials or scales (for example J.P. Morgan's asset-allocation views page), with a main asset-class dial and finer sub-asset dials. Should I record the main dials, the sub-asset dials, or both?

8 For full dial dashboards, should every dial become a row?

Some sources publish a full grid or dashboard with 30 or more asset dials. Do you want every dial recorded as its own row, or only the firm's highlighted or featured views? In the second test I recorded full grids faithfully, which was accurate but produced many rows from a single source.

9 Should a source ever include a firm's broader outlook material?

Right now I read exactly the document behind the provided link. Some reference answers appear to draw on a firm's broader publication set, such as separate bond, credit, or multi-asset pages. Should a source ever mean "this firm's outlook material generally"? If yes, please list the extra links that should be read for each source.

10 For a combined two-document row, which date and link should it carry?

When a single call is drawn from two documents read together (see question 2), the output row can only carry one date and one link. Which of the two should it use?

11 Is the reference I capture enough?

For each call I'm currently capturing both the page number and the exact supporting phrase. Web pages don't really have page numbers, so for those I just point to the exact phrase in the text. Is that enough of a reference for you, or would you want something more?

12 What confidence level should trigger mandatory human review?

The system flags lower-confidence calls for a human to check. Right now it flags anything below the top confidence band. Is that the right threshold, or would you like it set higher or lower?

Allocator Pro, proof of concept. Prepared for the Markets Recon team.