How to read AI engine outputs

A plain-language guide to the files each run produces.

Each run reads a set of fund and asset-manager outlook documents and produces two spreadsheets and one Word document. `output-client.csv` lists every allocation call we kept, one row per call. `failures-client.csv` lists everything we considered but did not keep, and why. The one-page firm summaries Word document gives a readable, per-firm digest of each manager's views. This guide explains the two spreadsheets and the actions to take.

output-client.csv - the calls we kept

Each row is one allocation call: a view on one sub-asset class, taken from one source document, with the supporting text. The first nine columns identify the call; the next four show the evidence behind it; the last three help you review it. The same file is supplied as `output-client.xlsx` for Excel.

Column	What it means
Firm	The asset manager that published the document.
Date	The document's own publication date, as DD/MM/YYYY. It is read from the document itself, not from any source list; it is left blank when the document does not state a date.
Source	The document title. When two related documents from the same firm are read together, their titles are joined with a `|`.
URL	A link to the document (also joined with a `|` for grouped documents).
Sub-Asset Class	The specific asset the call is about (e.g. "Emerging Market Equities"), taken from the approved asset list.
Asset Class Category	The broad grouping the sub-asset sits under. Filled in automatically from the asset list.
Canva Groupings	The display grouping for the sub-asset. Filled in automatically from the asset list.
Asset Class	The mid-level asset class for the sub-asset. Filled in automatically from the asset list.
View	The call: `O` overweight, `N` neutral, `U` underweight, `UNCERTAIN` when the document does not take a clear position.
Reasoning	Why the view was read that way, in one or two sentences.
Verbatim Evidence	The exact supporting text, copied from the document and nothing else, so you can check it against the source.

Locator	Where the evidence is: the page number for a PDF (plus the chart or table name when the call was read from one). For a document supplied as Markdown, a page is given only when the quote was found on exactly that one page of the matching PDF; otherwise it is left out and the Reviewer Notes say why. Blank for web pages and transcripts, which have no page numbers - search for the Verbatim Evidence instead.
Reviewer Notes	Notes about the call rather than from the document: what the second reader found, how a conflict between two passages was resolved, why the confidence was capped, and whether another document from the same firm holds a different view on the same asset.
confidence	A 0–100 score from fixed rules (evidence quality, how clearly the view is stated, whether a second reader confirmed it). It is *not* a model's own opinion of its confidence.
band	The confidence score in words: High, Medium, or Low.
review_flag	Whether a human should take a closer look. `strong_review` means the confidence score is below 50 (a Low band): check these first. `review` means a judgment was involved - two passages or documents disagreed and one was chosen, a second reader could not confirm the call, an implied reading challenges a stated one, or the view is `UNCERTAIN`. `none` means the call stands on its own; a Medium band on such a row only says the evidence was indirect or inferred, not that it needs a decision.

One call supported by several passages. When a firm states the same view in more than one document, or in more than one place, the call is kept once and every passage is preserved: the Reasoning, Verbatim Evidence and Locator cells then list one entry per document, each prefixed with the document title and separated by `||||`. **Two views on one asset.** The file holds one row per firm and asset. When two of a firm's documents take different positions on the same asset, the more recent document wins when both are dated on different days (otherwise: a view the document states over one derived from its forecasts, over an inferred one, then the higher confidence); the kept row is flagged, its Reviewer Notes name the other document and its view, and the other call is listed in `failures-client.csv` as "Superseded". Only when no rule can separate them are both rows kept, flagged, each naming the other.

failures-client.csv - what we considered but did not keep

This file records candidate calls that were set aside, and why. Most entries are healthy: duplicates that were kept once, near-duplicate calls, or views whose supporting text could not be verified and so were dropped rather than kept unconfirmed. A large failures file is normal and usually means the filtering is working. Rows are grouped by their **What happened** label and sorted with the most important labels first: entries that may need a decision or a re-run are at the top, and no-action housekeeping entries (duplicates, too-small changes) are at the bottom. The table below lists the labels in that same order:

What happened	Meaning
Document could not be ingested	A whole source couldn't be fetched or converted into text, so it was skipped while the rest of the run continued. Let us know so we can re-run it.
Conflicting views - none kept	Two views disagreed and could not be reconciled automatically; neither was kept. A human should decide.
Part of a document couldn't be read or processed / Second-reader check couldn't run / Skipped - source details missing	A processing issue affected part of a document. Let us know so we can re-run it.
Suggestion - review recommended	An implied reading challenges the firm's stated call; the stated call was kept and the challenge recorded for review.
Conflicting views - other call kept	The document took conflicting views on this asset; the call it holds most clearly was kept. Review if the kept call looks wrong.
Evidence could not be verified (or "… verified visually")	The supporting text couldn't be confirmed - even after a visual page check, for the "visually" variant - so the call was dropped rather than kept unverified. A human can check the cited page.
Evidence location missing	A chart/table call didn't record where it came from, so it couldn't be verified.
Evidence points the other way / Evidence is about a different asset / Not a forward-looking view	A second reader found the cited text doesn't back the call as made, so it was dropped. A human can check the cited page.
Skipped - asset not on the list	The asset named didn't match the approved list, so it was left out.
Change too small to call	The forecast change behind the call was too small to justify a view.
Evidence better supports a related asset	The same passage was used for more than one asset. It supports another asset more directly, so that call was kept and this one set aside; the notes name the kept asset. Review if this asset should also stand.
Superseded - a newer document from this firm updated this view	Another document from the same firm took a different view on this asset; a more current (or otherwise preferred) call was kept in its place. Review if the kept call looks wrong.
Merged into a related call - please review	The firm made views on two closely related assets (e.g. a broad category and a specific part of it) that pointed different ways; they were read as one collective call and the most relevant view was kept. Review that the kept view is the right one.

Duplicate - already covered	The same view for this asset appears in another of the firm's documents; kept once, with the other document's wording merged into the kept row's commentary. No action needed.
Duplicate - same call, related asset	Repeats a view already kept for a closely related asset from the same text; kept once. No action needed.
Merged - same view stated in several documents	Several of the firm's documents made the same call on this asset; they were merged into one row and every document's wording is preserved in that row's commentary. No action needed.
Merged - same view on a closely related asset	The firm made the same call on two closely related assets (e.g. a broad category and a specific part of it); they were merged into one row on the most precise asset and every document's wording is preserved. No action needed.

The **Explanation** and **Evidence / notes** columns in the file give the one-line meaning and the underlying text for each row. Because the file is sorted most-important-first, start at the top: act when you see "Document could not be ingested", "Conflicting views - none kept", "Suggestion - review recommended", or a repeated "Part of a document couldn't be…" for a document you expected calls from.

What to do with the output

- **Review the flagged rows first.** In `output-client.csv`, filter on every `review_flag` other than `none` (`strong_review` first, then `review`) - these are the calls we think most warrant a human eye.
- **Treat `UNCERTAIN` as a deliberate abstention, not an error.** It means the document did not take a clear position on that asset; we chose not to invent one.
- **Spot-check a call.** Read the Verbatim Evidence and confirm it against the Locator page in the source document (for a web page, search the page for the quoted text).
- **Send back anything that looks wrong.** If a kept call reads the wrong way, an expected call is missing, or a whole document produced no calls, let us know the Firm, Source, and Sub-Asset Class so we can look into it.